

Machine Learning Lab 6: Decision Tree Classification

Name: Deshmukh Pratik Bhushanrao

Roll No: 2448513

```
In [2]: import pandas as pd
import numpy as np
import seaborn as sns

#inferencing the essential libraries
```

```
In [3]: df = pd.read_excel('employee_data.xlsx')
df.head()

#importing the dataset
```

```
Out[3]:
```

	avg_monthly_hrs	department	filed_complaint	last_evaluation	n_projects	recently_p
0	221	engineering	NaN	0.932868	4	
1	232	support	NaN	NaN	3	
2	184	sales	NaN	0.788830	3	
3	206	sales	NaN	0.575688	4	
4	249	sales	NaN	0.845217	3	

1. Data Exploration and Visualization:

```
In [6]: df.info()

# checking for missing values
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14249 entries, 0 to 14248
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   avg_monthly_hrs        14249 non-null  int64
1   department             13540 non-null  object
2   filed_complaint        2058 non-null   float64
3   last_evaluation        12717 non-null  float64
4   n_projects             14249 non-null  int64
5   recently_promoted      300 non-null    float64
6   salary                 14249 non-null  object
7   satisfaction            14068 non-null  float64
8   status                 14249 non-null  object
9   tenure                 14068 non-null  float64
dtypes: float64(5), int64(2), object(3)
memory usage: 1.1+ MB
```

```
In [7]: df.isna().sum()
```

```
#checking for null values
```

```
Out[7]: avg_monthly_hrs      0
        department        709
        filed_complaint    12191
        last_evaluation     1532
        n_projects         0
        recently_promoted  13949
        salary             0
        satisfaction       181
        status             0
        tenure            181
        dtype: int64
```

```
In [8]: df.shape
```

```
#checking the shape of the dataset
```

```
Out[8]: (14249, 10)
```

```
In [10]: df.describe()
```

```
#checking the descriptive statistics
```

Out[10]:

	avg_monthly_hrs	filed_complaint	last_evaluation	n_projects	recently_promoted
count	14249.000000	2058.0	12717.000000	14249.000000	300
mean	199.795775	1.0	0.718477	3.773809	0.0
std	50.998714	0.0	0.173062	1.253126	0.0
min	49.000000	1.0	0.316175	1.000000	0.0
25%	155.000000	1.0	0.563866	3.000000	0.0
50%	199.000000	1.0	0.724939	4.000000	0.0
75%	245.000000	1.0	0.871358	5.000000	0.0
max	310.000000	1.0	1.000000	7.000000	0.0

In [11]: `df.drop(['filed_complaint', 'recently_promoted'], axis=1, inplace=True)`
`df.head()`

#dropping the unnecessary columns

Out[11]:

	avg_monthly_hrs	department	last_evaluation	n_projects	salary	satisfaction	status
0	221	engineering	0.932868	4	low	0.829896	
1	232	support	NaN	3	low	0.834544	Emp
2	184	sales	0.788830	3	medium	0.834988	Emp
3	206	sales	0.575688	4	low	0.424764	Emp
4	249	sales	0.845217	3	low	0.779043	Emp

In [12]: `df.isna().sum()`

#again checking for null values

Out[12]:

```
avg_monthly_hrs    0
department         709
last_evaluation    1532
n_projects         0
salary             0
satisfaction       181
status             0
tenure            181
dtype: int64
```

Data imputation

In [13]: `df['department'].fillna(df['department'].mode()[0], inplace=True)`
`df['last_evaluation'].fillna(df['last_evaluation'].mean(), inplace=True)`
`df['satisfaction'].fillna(df['satisfaction'].mean(), inplace=True)`
`df['tenure'].fillna(df['tenure'].mean(), inplace=True)`

#filling the missing values

C:\Users\prati\AppData\Local\Temp\ipykernel_1016\3877897122.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['department'].fillna(df['department'].mode()[0], inplace=True)
```

C:\Users\prati\AppData\Local\Temp\ipykernel_1016\3877897122.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['last_evaluation'].fillna(df['last_evaluation'].mean(), inplace=True)
```

C:\Users\prati\AppData\Local\Temp\ipykernel_1016\3877897122.py:3: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['satisfaction'].fillna(df['satisfaction'].mean(), inplace=True)
```

C:\Users\prati\AppData\Local\Temp\ipykernel_1016\3877897122.py:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['tenure'].fillna(df['tenure'].mean(), inplace=True)
```

In [14]: df.isna().sum()

#checking for null values


```
Out[14]: avg_monthly_hrs    0
         department        0
         last_evaluation    0
         n_projects        0
         salary            0
         satisfaction       0
         status            0
         tenure            0
         dtype: int64
```

```
In [15]: df.head()
```

#checking the dataset

```
Out[15]:
```

	avg_monthly_hrs	department	last_evaluation	n_projects	salary	satisfaction	s
0	221	engineering	0.932868	4	low	0.829896	
1	232	support	0.718477	3	low	0.834544	Emp
2	184	sales	0.788830	3	medium	0.834988	Emp
3	206	sales	0.575688	4	low	0.424764	Emp
4	249	sales	0.845217	3	low	0.779043	Emp



```
In [16]: df['department'].unique()
```

#checking for unique values


```
Out[16]: array(['engineering', 'support', 'sales', 'IT', 'product', 'marketing',
               'temp', 'procurement', 'finance', 'management',
               'information_technology', 'admin'], dtype=object)
```

visualizations

```
In [17]: sns.histplot(data=df, x='tenure', hue='status', bins=30, multiple='stack')
```


#checking the distribution of the target variable

```
Out[17]: <Axes: xlabel='tenure', ylabel='Count'>
```



```
In [18]: sns.boxplot(data=df,x='satisfaction',y='status')  
  
#checking for outliers
```

```
Out[18]: <Axes: xlabel='satisfaction', ylabel='status'>
```




```
In [20]: import matplotlib.pyplot as plt
```

```
In [21]: plt.pie(df["status"].value_counts(),labels=df["status"].unique(),autopct='%1.1f%%',
plt.plot()

#checking the distribution of the target variable
```


Out[21]: []



```
In [22]: sns.scatterplot(data=df[:50],x='satisfaction',y='last_evaluation',hue='status')

#checking for correlation
```

Out[22]: <Axes: xlabel='satisfaction', ylabel='last_evaluation'>



Encoding

```
In [24]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df['department'] = le.fit_transform(df['department'])
df['salary'] = le.fit_transform(df['salary'])
df['status'] = le.fit_transform(df['status'])
df.head()

#encoding the categorical variables
```

```
Out[24]:
```

	avg_monthly_hrs	department	last_evaluation	n_projects	salary	satisfaction	status
0	221	2	0.932868	4	1	0.829896	1
1	232	10	0.718477	3	1	0.834544	0
2	184	9	0.788830	3	2	0.834988	0
3	206	9	0.575688	4	1	0.424764	0
4	249	9	0.845217	3	1	0.779043	0

```
In [25]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
x = pd.DataFrame(sc.fit_transform(df.drop('status', axis=1)), columns=df.columns)
df_scaled = pd.concat([x, df['status']], axis=1)
df_scaled.head()

#standardizing(scaling) the dataset
```


Out[25]:


	avg_monthly_hrs	department	last_evaluation	n_projects	salary	satisfaction	
0	0.415794	-1.205166	1.311357	0.180508	-0.561893	0.838210	1.0
1	0.631493	1.055552	0.000000	-0.617524	-0.561893	0.856885	-1.0
2	-0.309740	0.772962	0.430327	-0.617524	1.040992	0.858668	-0.0
3	0.121659	0.772962	-0.873391	0.180508	-0.561893	-0.789711	-1.0
4	0.964847	0.772962	0.775231	-0.617524	-0.561893	0.633869	-0.0

In [26]:

```
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 10))
sns.heatmap(df_scaled.corr(), annot=True)
plt.show()

#checking for correlation
```



In []:

2. Decision Tree Modeling with Tree Pruning and Split Criteria:

Split the dataset into a training set and a testing set (e.g., 80% training, 20% testing).

```
In [27]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(df_scaled.drop('status', axis=1), df_scaled['status'],
#splitting the dataset into train and test sets with 80:20 ratio
```

Build a decision tree classifier to predict employee attrition based on selected features (e.g., tenure, satisfaction, number of projects).

```
In [28]: from sklearn.tree import DecisionTreeClassifier
cnf_gini = DecisionTreeClassifier(criterion='gini')
cnf_gini.fit(X_train, y_train)
y_pred = cnf_gini.predict(X_test)

cnf_gini.score(X_test, y_test)

#checking building the model with Gini impurity criterion and checking the accuracy
```

Out[28]: 0.9578947368421052

```
In [29]: cnf_entropy = DecisionTreeClassifier(criterion='entropy')
cnf_entropy.fit(X_train, y_train)
y_pred = cnf_entropy.predict(X_test)

cnf_entropy.score(X_test, y_test)

#checking building the model with Entropy criterion and checking the accuracy
```

Out[29]: 0.9652631578947368

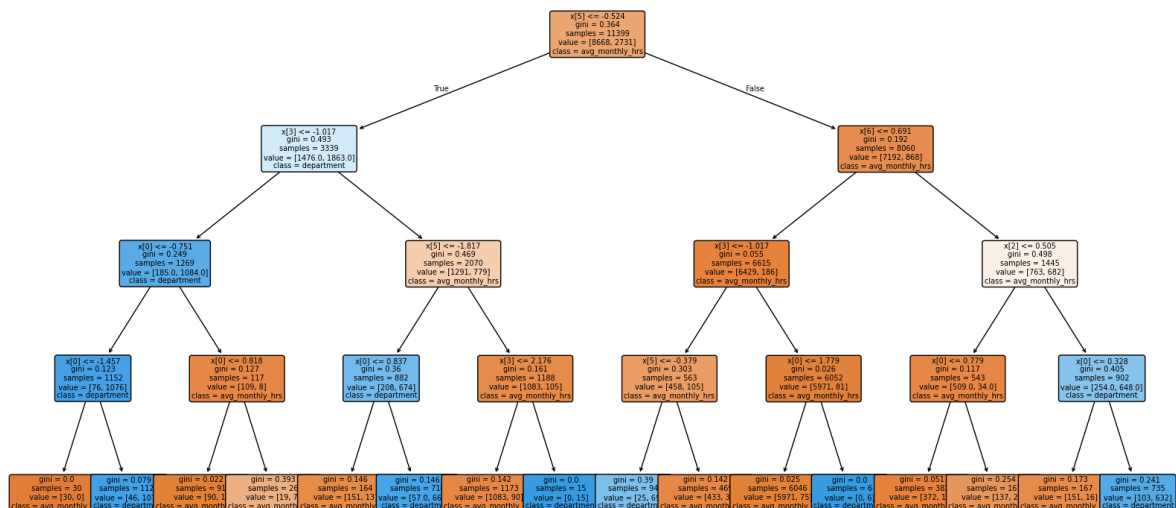
Visualize the decision tree structure. How deep is the tree, and what are the most influential features for predicting attrition?

```
In [30]: from sklearn.tree import export_text, plot_tree

example_tree = DecisionTreeClassifier(criterion='gini', max_depth=4, min_samples_leaf=10)
plt.figure(figsize=(20, 10))
plot_tree(example_tree, class_names=X_train.columns, filled=True, rounded=True,
plt.plot())

#plotting the decision tree
```

Out[30]: []



Apply post-pruning techniques to control the complexity of the tree and prevent overfitting. Experiment with different pruning strategies, such as minimum leaf size or maximum depth, to find the optimal tree size

```
In [31]: path = cnf_gini.cost_complexity_pruning_path(X_train, y_train)
ccp_alphas, impurities = path.ccp_alphas, path.impurities

pruned_models = []

for ccp_alphas in ccp_alphas:
    pruned_model = DecisionTreeClassifier(criterion='gini', ccp_alpha=ccp_alphas)
    pruned_models.append(pruned_model)

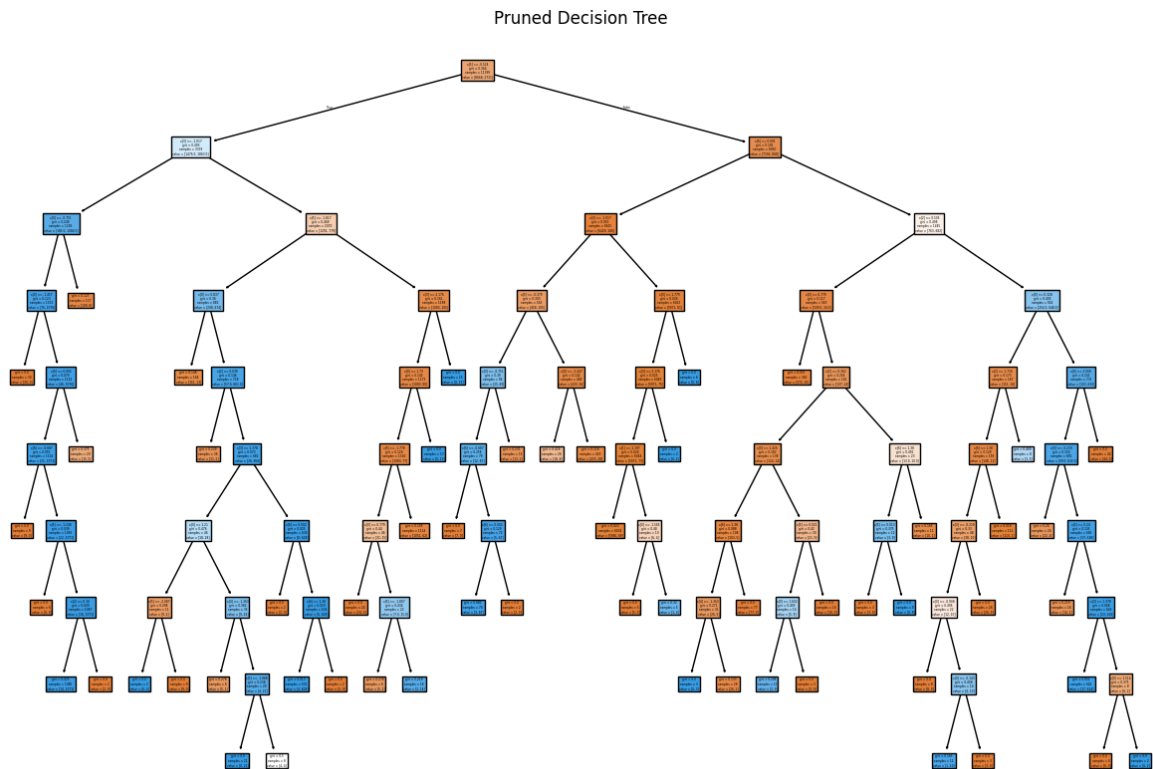
best_accuracy = 0
best_pruned_model = None

for pruned_model in pruned_models:
    accuracy = pruned_model.score(X_test, y_test)
    if accuracy > best_accuracy:
        best_accuracy = accuracy
        best_pruned_model = pruned_model

# finding the best pruned model based on accuracy
```

```
In [32]: plt.figure(figsize=(15, 10))
plot_tree(best_pruned_model, filled=True)
plt.title("Pruned Decision Tree")
plt.show()

# plotting the pruned decision tree
```



```
In [33]: from sklearn.model_selection import GridSearchCV
model = DecisionTreeClassifier()
params = {
    'criterion': ['gini', 'entropy'],
    'max_depth': list(np.random.randint(2, 20, 1)),
    'min_samples_leaf': [1, 2, 4, 6]
}
search = GridSearchCV(model, params, scoring='accuracy').fit(X_train, y_train)

#finding the best model
```

```
In [34]: print(search.best_params_)
print(search.best_score_)

#finding the best parameters and best score

{'criterion': 'gini', 'max_depth': 4, 'min_samples_leaf': 1}
0.9535926037120005
```

```
In [35]: bestModel = search.best_estimator_
y_pred = bestModel.predict(X_test)
bestModel.score(X_test, y_test)

#checking the accuracy of the best model
```

Out[35]: 0.9624561403508772

3. Model Evaluation

```
In [37]: from sklearn.metrics import confusion_matrix, classification_report
```

```
In [38]: conf_mat = confusion_matrix(y_test, y_pred)
print(conf_mat, end="\n\n")
print(classification_report(y_test, y_pred))

#checking the confusion matrix and classification report
```

```
[[2135   54]
 [  53  608]]
```

	precision	recall	f1-score	support
0	0.98	0.98	0.98	2189
1	0.92	0.92	0.92	661
accuracy			0.96	2850
macro avg	0.95	0.95	0.95	2850
weighted avg	0.96	0.96	0.96	2850

4. Feature Importance Visualization:

```
In [40]: tn, fp, fn, tp = conf_mat.ravel()
labels = ['True Positives', 'True Negatives', 'False Positives', 'False Negatives']


sns.barplot(y=labels, x=[tp, tn, fp, fn], palette=['green', 'green', 'red', 'red'],
plt.show()

#plotting the confusion matrix with barchart
```

C:\Users\prati\AppData\Local\Temp\ipykernel_1016\342439456.py:4: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(y=labels, x=[tp, tn, fp, fn], palette=['green', 'green', 'red', 'red'],
legend=False)
```



```
In [41]: entropy_imp = cnf_entropy.feature_importances_
gini_imp = cnf_gini.feature_importances_
features = X_test.columns

# Plotting the bar chart
fig, ax = plt.subplots(figsize=(15, 10))
bar_width = 0.35
index = np.arange(len(features))


# Gini bar chart
bar1 = plt.bar(index, gini_imp, bar_width, label='Gini', color='r')

# Entropy bar chart
bar2 = plt.bar(index + bar_width, entropy_imp, bar_width, label='Entropy', co

# Adding Labels and Titles
plt.xlabel('Features')
plt.ylabel('Importance')
plt.title('Feature Importance in Pruned Decision Tree (Gini vs Entropy)')
plt.xticks(index + bar_width / 2, features, rotation=45)
plt.legend()

plt.show()

#plotting the feature importance after pruning
```



-----EOD-----